LLM / VLM References

LLM 基础、Scaling 与架构

- Vaswani et al., 2017, Attention Is All You Need
- Brown et al., 2020, Language Models are Few-Shot Learners
- Kaplan et al., 2020, Scaling Laws for Neural Language Models
- Hoffmann et al., 2022, Training Compute-Optimal Large Language Models
- Dao et al., 2022, FlashAttention
- Meta AI, 2024, The Llama 3 Herd of Models
- Meta AI, 2025, Llama 4 multimodal intelligence
- Google, 2025, Gemini 2.5 Technical Report
- Qwen Team, 2024, Qwen2.5 Technical Report
- Qwen Team, 2025, Qwen3 Technical Report

Post-training、Alignment 与 Reasoning

- Ouyang et al., 2022, Training language models to follow instructions with human feedback
- Bai et al., 2022, Constitutional AI: Harmlessness from AI Feedback
- Lightman et al., 2023, Let's Verify Step by Step
- Rafailov et al., 2023, Direct Preference Optimization
- Shao et al., 2024, DeepSeekMath: Pushing the Limits of Mathematical Reasoning in Open Language Models
- Guan et al., 2024, Deliberative Alignment: Reasoning Enables Safer Language Models
- DeepSeek-AI, 2025, DeepSeek-R1: Incentivizing Reasoning Capability in LLMs via Reinforcement Learning
- OpenAI, 2025, OpenAI o3 and o4-mini System Card
- OpenAI, 2025, Introducing OpenAI o3 and o4-mini
- OpenAI, 2026, OpenAI GPT-5 System Card
- Gould et al., 2026, Think Fast: Estimating No-CoT Task-Completion Time Horizons of Frontier AI Models

PEFT 与高效微调

- Hu et al., 2021, LoRA
- Dettmers et al., 2023, QLoRA

Prompt、RAG、Agent 与 Eval

- Wei et al., 2022, Chain-of-Thought Prompting

- Wang et al., 2022, Self-Consistency Improves Chain of Thought Reasoning
- Lewis et al., 2020, Retrieval-Augmented Generation
- Guu et al., 2020, REALM
- Microsoft, GraphRAG
- VibrantLabs, RAGAS
- Singh et al., 2025, Agentic Retrieval-Augmented Generation
- Sharma, 2025, Retrieval-Augmented Generation: A Comprehensive Survey
- Schick et al., 2023, Toolformer
- Yao et al., 2022, ReAct
- Liu et al., 2023, AgentBench
- Jimenez et al., 2023, SWE-bench
- Yao et al., 2024, TAU-bench
- Wei et al., 2025, BrowseComp
- Kapoor et al., 2026, Open-World Evaluations for Measuring Frontier AI Capabilities
- Liang et al., 2022, HELM
- EleutherAI, LM Evaluation Harness
- OpenAI, Evals
- Lee et al., 2024, VHELM
- Yao et al., 2026, Harness-Bench
- Muennighoff et al., 2022, MTEB
- Hendrycks et al., 2020, MMLU

2025-2026 SOTA 技术报告和官方模型页

- OpenAI, 2025, GPT-5 is here
- OpenAI, 2026, Introducing GPT-5.5
- Anthropic, 2025, Claude 4
- Anthropic, 2026, Claude Opus 4.7
- Google, 2025, Gemini 2.5 Technical Report
- DeepSeek-AI, 2025, DeepSeek-R1
- Qwen Team, 2025, Qwen3 Technical Report
- Kimi Team, 2025, Kimi K2: Open Agentic Intelligence
- Kimi Team, 2026, Kimi K2.5: Visual Agentic Intelligence
- GLM Team, 2026, GLM-5: from Vibe Coding to Agentic Engineering
- Z.ai, 2026, GLM-5.2 / GLM-5.1 / GLM-5 official repo
- GLM Team, 2025, GLM-4.5: Agentic, Reasoning, and Coding Foundation Models
- GLM-V Team, 2025, GLM-4.1V-Thinking and GLM-4.5V

VLM / MLLM

- Dosovitskiy et al., 2020, Vision Transformer
- Radford et al., 2021, CLIP
- Alayrac et al., 2022, Flamingo

- Li et al., 2023, BLIP-2
- Liu et al., 2023, LLaVA: Visual Instruction Tuning
- Bai et al., 2025, Qwen2.5-VL Technical Report
- Bai et al., 2025, Qwen2.5 Technical Report
- OpenAI, 2023, GPT-4 Technical Report
- Liu et al., 2023, MMBench
- Yue et al., 2023, MMMU

章节索引 ← 上一章 下一章 →