

LLM / VLM References

R.1 LLM 基础、Scaling 与架构

- Vaswani et al., 2017, [Attention Is All You Need](#)
- Brown et al., 2020, [Language Models are Few-Shot Learners](#)
- Kaplan et al., 2020, [Scaling Laws for Neural Language Models](#)
- Hoffmann et al., 2022, [Training Compute-Optimal Large Language Models](#)
- Dao et al., 2022, [FlashAttention](#)
- Meta AI, 2024, [The Llama 3 Herd of Models](#)
- Meta AI, 2025, [Llama 4 multimodal intelligence](#)
- Google, 2025, [Gemini 2.5 Technical Report](#)
- Qwen Team, 2024, [Qwen2.5 Technical Report](#)
- Qwen Team, 2025, [Qwen3 Technical Report](#)

R.2 Post-training、Alignment 与 Reasoning

- Ouyang et al., 2022, [Training language models to follow instructions with human feedback](#)
- Bai et al., 2022, [Constitutional AI: Harmlessness from AI Feedback](#)
- Lightman et al., 2023, [Let's Verify Step by Step](#)
- Rafailov et al., 2023, [Direct Preference Optimization](#)
- Shao et al., 2024, [DeepSeekMath: Pushing the Limits of Mathematical Reasoning in Open Language Models](#)
- Guan et al., 2024, [Deliberative Alignment: Reasoning Enables Safer Language Models](#)
- DeepSeek-AI, 2025, [DeepSeek-R1: Incentivizing Reasoning Capability in LLMs via Reinforcement Learning](#)
- OpenAI, 2025, [OpenAI o3 and o4-mini System Card](#)
- OpenAI, 2025, [Introducing OpenAI o3 and o4-mini](#)
- OpenAI, 2026, [OpenAI GPT-5 System Card](#)
- Gould et al., 2026, [Think Fast: Estimating No-CoT Task-Completion Time Horizons of Frontier AI Models](#)

R.3 PEFT 与高效微调

- Hu et al., 2021, [LoRA](#)
- Dettmers et al., 2023, [QLoRA](#)

R.4 Prompt、RAG、Agent 与 Eval

- Wei et al., 2022, [Chain-of-Thought Prompting](#)

- Wang et al., 2022, [Self-Consistency Improves Chain of Thought Reasoning](#)
- Lewis et al., 2020, [Retrieval-Augmented Generation](#)
- Guu et al., 2020, [REALM](#)
- Microsoft, [GraphRAG](#)
- VibrantLabs, [RAGAS](#)
- Singh et al., 2025, [Agentic Retrieval-Augmented Generation](#)
- Sharma, 2025, [Retrieval-Augmented Generation: A Comprehensive Survey](#)
- Schick et al., 2023, [Toolformer](#)
- Yao et al., 2022, [ReAct](#)
- Liu et al., 2023, [AgentBench](#)
- Jimenez et al., 2023, [SWE-bench](#)
- Yao et al., 2024, [TAU-bench](#)
- Wei et al., 2025, [BrowseComp](#)
- Kapoor et al., 2026, [Open-World Evaluations for Measuring Frontier AI Capabilities](#)
- Liang et al., 2022, [HELM](#)
- EleutherAI, [LM Evaluation Harness](#)
- OpenAI, [Evals](#)
- Lee et al., 2024, [VHELM](#)
- Yao et al., 2026, [Harness-Bench](#)
- Muennighoff et al., 2022, [MTEB](#)
- Hendrycks et al., 2020, [MMLU](#)

R.5 2025-2026 SOTA 技术报告和官方模型页

- OpenAI, 2025, [GPT-5 is here](#)
- OpenAI, 2026, [Introducing GPT-5.5](#)
- Anthropic, 2025, [Claude 4](#)
- Anthropic, 2026, [Claude Opus 4.7](#)
- Google, 2025, [Gemini 2.5 Technical Report](#)
- DeepSeek-AI, 2025, [DeepSeek-R1](#)
- Qwen Team, 2025, [Qwen3 Technical Report](#)
- Kimi Team, 2025, [Kimi K2: Open Agentic Intelligence](#)
- Kimi Team, 2026, [Kimi K2.5: Visual Agentic Intelligence](#)
- GLM Team, 2026, [GLM-5: from Vibe Coding to Agentic Engineering](#)
- Z.ai, 2026, [GLM-5.2 / GLM-5.1 / GLM-5 official repo](#)
- GLM Team, 2025, [GLM-4.5: Agentic, Reasoning, and Coding Foundation Models](#)
- GLM-V Team, 2025, [GLM-4.1V-Thinking and GLM-4.5V](#)

R.6 VLM / MLLM

- Dosovitskiy et al., 2020, [Vision Transformer](#)
- Radford et al., 2021, [CLIP](#)
- Alayrac et al., 2022, [Flamingo](#)

- Li et al., 2023, [BLIP-2](#)
- Liu et al., 2023, [LLaVA: Visual Instruction Tuning](#)
- Bai et al., 2025, [Qwen2.5-VL Technical Report](#)
- Bai et al., 2025, [Qwen2.5 Technical Report](#)
- OpenAI, 2023, [GPT-4 Technical Report](#)
- Liu et al., 2023, [MMBench](#)
- Yue et al., 2023, [MMMU](#)

[章节索引](#)

[← 上一章](#) [下一章 →](#)
